

Introduction to Operating Systems

Introduction to Operating Systems

Overview of Operating System Concepts

The primary goal of an Operating System (OS) is to manage computer hardware resources and provide a platform for running application software. **Operating System** is a program that manages computer hardware and software resources, providing a platform for running application software.

Operating System Objectives

The main objectives of an Operating System are:

1. **Process Management:** to manage multiple programs running concurrently
2. **Memory Management:** to manage memory allocation and deallocation
3. **Storage Management:** to manage file systems and storage devices
4. **Protection and Security:** to provide mechanisms for controlling access to computer resources
5. **Input/Output Management:** to manage input/output operations between devices and programs

User View and System View

There are two views of an Operating System:

- **User View:** the Operating System is seen as a program that manages computer hardware and provides services to applications
- **System View:** the Operating System is seen as a resource manager that allocates and deallocates resources such as CPU, memory, and I/O devices

Operating System Definition

An **Operating System** is a program that:

- Manages computer hardware resources
- Provides a platform for running application software
- Acts as an intermediary between computer hardware and application software

Computer System Architecture

A computer system consists of:

- **Hardware:** CPU, memory, input/output devices
- **Software:** Operating System, application software
- **Firmware:** permanent software stored in non-volatile memory

Components of Computer System Architecture

- **Central Processing Unit (CPU):** executes instructions and performs calculations
- **Memory:** stores data and program instructions
- **Input/Output Devices:** allow users to interact with the computer system

OS Structure and Operations

The Operating System is divided into two main parts:

- **Kernel:** the core part of the Operating System that manages hardware resources
- **System Programs:** provide services to applications and users

Operating System Structure

- **Monolithic Structure:** the Operating System is a single, self-contained program
- **Modular Structure:** the Operating System is divided into separate modules or components

Process Management

Process is a program in execution, and **process management** is the ability of the Operating System to create, execute, and terminate processes.

Process States

- **Running:** the process is currently executing
- **Waiting:** the process is waiting for a resource or event to occur
- **Blocked:** the process is blocked and cannot execute

Memory Management

Memory management is the ability of the Operating System to manage memory allocation and deallocation.

Memory Allocation

- **Static Allocation:** memory is allocated at compile time
- **Dynamic Allocation:** memory is allocated at runtime

Storage Management

Storage management is the ability of the Operating System to manage file systems and storage devices.

File Systems

- **File:** a collection of related data
- **File System:** a system for storing and retrieving files

Protection and Security

Protection is the ability of the Operating System to control access to computer resources, and **security** is the ability of the Operating System to prevent unauthorized access to computer resources.

Access Control Mechanisms

- **Access Control Lists (ACLs):** a list of users and their access rights
- **Capabilities:** a token that grants access to a resource

Computing Environments

Computing environments refer to the different types of computing systems, including:

- **Desktop Computing:** a single user system
- **Mainframe Computing:** a large-scale system for multiple users
- **Embedded Systems:** a system that is embedded in a device or appliance

Operating System Services

The Operating System provides various services, including:

- **Process Creation:** creating a new process
- **Memory Allocation:** allocating memory to a process
- **Input/Output Management:** managing input/output operations

User and OS Interface

The **user interface** is the interface between the user and the Operating System, and the **OS interface** is the interface between the Operating System and application software.

Command-Line Interface

- **Commands:** instructions to the Operating System
- **Arguments:** parameters passed to a command

System Calls and Types

A **system call** is a request to the Operating System to perform a service, and there are different types of system calls, including:

- **Process Control System Calls:** create, execute, and terminate processes
- **File System System Calls:** manage file systems and storage devices

System Programs and OS Design

System programs provide services to applications and users, and **OS design** refers to the design of the Operating System, including its structure and components.

System Programs

- **Compiler:** translates source code into object code
- **Loader:** loads object code into memory

OS Structure and Implementation

The **OS structure** refers to the organization of the Operating System, and the **implementation** refers to the actual coding of the Operating System.

OS Implementation

- **Kernel Implementation:** the implementation of the kernel
- **System Program Implementation:** the implementation of system programs

CPU Scheduling

CPU Scheduling

Introduction to Process Concepts

The concept of a **process** refers to a program in execution, including the current activity, memory, and system resources. A process has several key characteristics:

- **Process ID (PID):** a unique identifier for each process
- **Program Counter:** a register that stores the address of the next instruction to be executed
- **Memory:** the process's allocated memory space
- **Open Files:** files currently being accessed by the process

Process State

A process can be in one of several states:

1. **Running:** the process is currently executing
2. **Ready:** the process is waiting to be executed
3. **Waiting:** the process is waiting for an event to occur
4. **Zombie:** the process has finished execution but still has an entry in the process table
5. **Dead:** the process has finished execution and has been removed from the process table

Process Control Block

A **Process Control Block (PCB)** is a data structure that contains information about a process, including:

- **Process ID**
- **Program Counter**
- **Memory:** allocated memory space
- **Open Files:** files currently being accessed
- **Scheduling Information:** priority, scheduling algorithm, etc. The PCB is used by the operating system to manage and schedule processes.

Threads and Process Scheduling

Threads are lightweight processes that share the same memory space and resources. Thread scheduling is the process of allocating CPU time to threads.

Types of Scheduling

1. **Preemptive Scheduling:** the operating system can interrupt a process and allocate the CPU to another process

2. **Non-Preemptive Scheduling:** the operating system cannot interrupt a process and must wait for it to finish execution

Scheduling Queues and Schedulers

A **scheduling queue** is a data structure that stores processes waiting to be executed. A **scheduler** is a program that allocates CPU time to processes.

Types of Schedulers

1. **Long-Term Scheduler:** selects which processes to load into memory
2. **Short-Term Scheduler:** selects which process to execute next
3. **Medium-Term Scheduler:** selects which process to swap out of memory

Context Switch and Operations on Processes

A **context switch** is the process of switching between two processes. Operations on processes include:

1. **Create:** create a new process
2. **Delete:** delete a process
3. **Suspend:** suspend a process
4. **Resume:** resume a suspended process

System Calls for Process Management

System calls are APIs that allow processes to interact with the operating system. System calls for process management include:

1. **Fork:** create a new process
2. **Exec:** execute a new program
3. **Wait:** wait for a process to finish execution
4. **Kill:** terminate a process

Inter-Process Communication

Inter-Process Communication (IPC) is the exchange of data between processes. Types of IPC include:

1. **Pipes:** a one-way communication channel
2. **Named Pipes:** a named pipe that can be accessed by multiple processes
3. **Message Queues:** a queue that stores messages between processes
4. **Shared Memory:** a shared memory space that can be accessed by multiple processes

Ordinary Pipes and Named Pipes

Ordinary Pipes are a one-way communication channel between two processes. **Named Pipes** are a named pipe that can be accessed by multiple processes.

Ordinary Pipes

1. **Pipe:** a pipe is created using the `pipe()` system call
2. **Read:** a process can read from the pipe using the `read()` system call
3. **Write:** a process can write to the pipe using the `write()` system call

Message Queues and Shared Memory

Message Queues are a queue that stores messages between processes. **Shared Memory** is a shared memory space that can be accessed by multiple processes.

Message Queues

1. **Create:** a message queue is created using the `msgget()` system call
2. **Send:** a process can send a message to the queue using the `msgsnd()` system call
3. **Receive:** a process can receive a message from the queue using the `msgrcv()` system call

Scheduling Criteria and Algorithms

Scheduling criteria include:

1. **CPU Utilization:** the percentage of time the CPU is busy
 2. **Throughput:** the number of processes completed per unit time
 3. **Turnaround Time:** the time it takes for a process to complete
 4. **Waiting Time:** the time a process waits in the ready queue
- Scheduling algorithms** include:
5. **First-Come-First-Served (FCFS):** the process that arrives first is executed first
 6. **Shortest Job First (SJF):** the process with the shortest execution time is executed first
 7. **Priority Scheduling:** the process with the highest priority is executed first

Multiple Processor Scheduling

Multiple processor scheduling is the process of scheduling processes on multiple processors.

Types of Multiple Processor Scheduling

1. **Asymmetric Multiprocessing:** one processor is the master and the others are slaves
2. **Symmetric Multiprocessing:** all processors are equal and can execute any process

Real-Time Scheduling

Real-time scheduling is the process of scheduling processes with strict deadlines.

Types of Real-Time Scheduling

1. **Hard Real-Time Scheduling**: the deadline must be met
2. **Soft Real-Time Scheduling**: the deadline is desirable but not required

Thread Scheduling

Thread scheduling is the process of scheduling threads.

Types of Thread Scheduling

1. **Preemptive Thread Scheduling**: the operating system can interrupt a thread and allocate the CPU to another thread
2. **Non-Preemptive Thread Scheduling**: the operating system cannot interrupt a thread and must wait for it to finish execution

Linux Scheduling

Linux scheduling uses a combination of scheduling algorithms, including:

1. **CFS (Completely Fair Scheduler)**: a scheduling algorithm that allocates CPU time fairly among processes
2. **O(1) Scheduler**: a scheduling algorithm that schedules processes in $O(1)$ time

Windows Scheduling

Windows scheduling uses a combination of scheduling algorithms, including:

1. **Multilevel Feedback Queue**: a scheduling algorithm that uses multiple queues with different priorities
2. **Rate Monotonic Scheduling**: a scheduling algorithm that schedules processes based on their priority and execution time

Process Synchronization and Deadlocks

Process Synchronization and Deadlocks

Introduction to Process Synchronization

Process synchronization refers to the coordination of multiple processes or threads that share common resources to prevent conflicts and ensure data consistency. **Process synchronization** is necessary to prevent problems such as **race conditions**, **deadlocks**, and **starvation**.

Background and Critical Section Problem

The critical section problem is a classic problem in process synchronization where multiple processes share a common resource, and each process has a **critical section** where it accesses the shared resource. To solve this problem, we need to ensure that only one process can execute its critical section at a time.

Peterson's Solution

Peterson's solution is a software-based solution to the critical section problem. It uses a combination of **flags** and **turns** to ensure mutual exclusion. The algorithm works as follows:

1. Each process sets its own flag to indicate that it is ready to enter its critical section.
2. If the other process's flag is not set, the process can enter its critical section.
3. If the other process's flag is set, the process sets its own turn variable to indicate that it is waiting for its turn.
4. The process then waits for its turn to enter its critical section.

Synchronization Hardware

Synchronization hardware provides a way to implement process synchronization using hardware components. The most common synchronization hardware components are:

- **Test-and-set** instructions: These instructions test the value of a variable and set it to a new value in a single operation.
- **Swap** instructions: These instructions swap the values of two variables in a single operation.
- **Locks**: These are hardware components that can be used to implement mutual exclusion.

Semaphores

A **semaphore** is a variable that controls access to a shared resource. Semaphores can be used to implement mutual exclusion, synchronization, and communication between processes. There are two types of semaphores:

- **Binary semaphores**: These semaphores can have only two values: 0 and 1.
- **Counting semaphores**: These semaphores can have any non-negative integer value.

Semaphores Operations

Semaphores support two operations:

- **Wait:** This operation decrements the value of the semaphore. If the value is 0, the process is blocked until the semaphore is signaled.
- **Signal:** This operation increments the value of the semaphore. If there are any processes waiting on the semaphore, one of them is unblocked.

Classic Problems of Synchronization

The classic problems of synchronization are:

- **Dining Philosophers Problem:** This problem involves five philosophers who share a common table and need to pick up two chopsticks to eat.
- **Readers-Writers Problem:** This problem involves a shared resource that can be accessed by multiple readers and writers.
- **Sleeping Barber Problem:** This problem involves a barber who needs to synchronize with multiple customers.

Monitors

A **monitor** is a high-level synchronization construct that provides a way to implement mutual exclusion and synchronization. A monitor consists of:

- **Condition variables:** These variables are used to signal between processes.
- **Locks:** These are used to implement mutual exclusion.

Synchronization in Linux and Windows

Both Linux and Windows provide synchronization primitives such as **mutexes**, **semaphores**, and **condition variables**. These primitives can be used to implement process synchronization and communication.

System Model for Deadlocks

A **deadlock** occurs when two or more processes are blocked indefinitely, each waiting for the other to release a resource. The system model for deadlocks consists of:

- **Resources:** These are the shared resources that can be accessed by multiple processes.
- **Processes:** These are the entities that compete for the shared resources.
- **Resource allocation:** This refers to the allocation of resources to processes.

Deadlock Characterization

Deadlocks can be characterized by the following properties:

- **Mutual exclusion:** Each resource is either available or allocated to a single process.
- **Hold and wait:** A process is holding a resource and waiting for another resource.
- **No preemption:** A process cannot be preempted while holding a resource.
- **Circular wait:** A process is waiting for a resource that is held by another process, which is waiting for a resource held by the first process.

Methods for Handling Deadlocks

There are three methods for handling deadlocks:

1. **Deadlock prevention:** This involves preventing deadlocks from occurring by ensuring that the deadlock properties are not met.
2. **Deadlock avoidance:** This involves avoiding deadlocks by carefully allocating resources to processes.
3. **Deadlock detection and recovery:** This involves detecting deadlocks and recovering from them.

Deadlock Prevention

Deadlock prevention involves ensuring that the deadlock properties are not met. This can be done by:

- **Eliminating mutual exclusion:** By allowing multiple processes to access a resource simultaneously.
- **Eliminating hold and wait:** By ensuring that a process does not hold a resource while waiting for another resource.
- **Eliminating no preemption:** By allowing a process to be preempted while holding a resource.
- **Eliminating circular wait:** By ordering the resources in a way that prevents circular wait.

Deadlock Avoidance

Deadlock avoidance involves carefully allocating resources to processes to avoid deadlocks. This can be done using:

- **Banker's algorithm:** This algorithm ensures that the system is always in a safe state, where it is possible to allocate resources to processes without causing a deadlock.
- **Resource ordering:** This involves ordering the resources in a way that prevents deadlocks.

Deadlock Detection and Recovery

Deadlock detection involves detecting deadlocks in the system. This can be done using:

- **Wait-for graphs:** These graphs represent the wait-for relationships between processes.
- **Cycle detection:** This involves detecting cycles in the wait-for graph, which indicate a deadlock.

Recovery from deadlock involves:

- **Process termination:** This involves terminating one or more processes to break the deadlock.
- **Resource preemption:** This involves preempting a resource from a process to break the deadlock.

Case Studies of Deadlock Handling

Deadlock handling is critical in many systems, including:

- **Databases:** Deadlocks can occur in databases when multiple transactions access the same resources.
- **Operating systems:** Deadlocks can occur in operating systems when multiple processes access the same resources.
- **Networks:** Deadlocks can occur in networks when multiple nodes access the same resources.

Memory Management

Memory Management

Memory management refers to the process of managing the memory resources of a computer system. It is a critical component of operating system design, as it enables multiple programs to share the same memory space efficiently. **Memory management** involves allocating memory to programs, deallocating memory when it is no longer needed, and preventing programs from accessing memory that has not been allocated to them.

Background and Memory Management

Memory management is essential because it allows multiple programs to run simultaneously, sharing the same memory space. Without memory management, programs would have to be

run one at a time, and the system would be unable to efficiently utilize its memory resources.

Memory management techniques include:

- **Memory allocation:** The process of assigning memory to a program or process.
- **Memory deallocation:** The process of freeing up memory that is no longer needed by a program or process.
- **Memory protection:** The process of preventing programs from accessing memory that has not been allocated to them.

Swapping and Contiguous Memory Allocation

Swapping is a memory management technique that involves moving a program's memory from RAM to disk storage when the program is not actively running. **Contiguous memory allocation** is a technique that involves allocating a single block of memory to a program. The advantages of contiguous memory allocation include:

- **Efficient memory use:** Contiguous memory allocation can reduce memory fragmentation, which occurs when free memory is broken into small, non-contiguous blocks.
- **Fast memory access:** Contiguous memory allocation can improve memory access times, as the system does not have to search for non-contiguous blocks of memory.

Segmentation and Paging

Segmentation is a memory management technique that involves dividing a program's memory into smaller segments, each with its own set of permissions. **Paging** is a technique that involves dividing a program's memory into small, fixed-size blocks called pages. The advantages of segmentation and paging include:

- **Improved memory protection:** Segmentation and paging can improve memory protection by preventing programs from accessing memory that has not been allocated to them.
- **Efficient memory use:** Segmentation and paging can reduce memory fragmentation, which occurs when free memory is broken into small, non-contiguous blocks.

Structure of Page Table

A **page table** is a data structure that maps virtual page numbers to physical page numbers. The page table contains the following components:

- **Page number:** The virtual page number of the page.
- **Frame number:** The physical frame number of the page.
- **Valid bit:** A bit that indicates whether the page is valid or not.

- **Dirty bit:** A bit that indicates whether the page has been modified or not.

Virtual Memory Management

Virtual memory is a memory management technique that involves using disk storage to augment the main memory of a computer system. Virtual memory management involves:

- **Page faults:** When a program accesses a page that is not in main memory, a page fault occurs, and the system must retrieve the page from disk storage.
- **Page replacement:** When a page fault occurs, the system must replace an existing page in main memory with the new page.

Demand Paging

Demand paging is a virtual memory management technique that involves loading pages into main memory only when they are needed. The advantages of demand paging include:

- **Improved memory use:** Demand paging can reduce the amount of memory required by a program, as only the pages that are actually needed are loaded into main memory.
- **Fast program startup:** Demand paging can improve program startup times, as the system does not have to load all of the program's pages into main memory at once.

Copy-on-Write and Page Replacement

Copy-on-write is a technique that involves creating a copy of a page when a program attempts to modify it. **Page replacement algorithms** are used to determine which page to replace when a page fault occurs. Common page replacement algorithms include:

- **First-in, first-out (FIFO):** The page that has been in main memory the longest is replaced.
- **Least recently used (LRU):** The page that has not been accessed for the longest period of time is replaced.

Allocation of Frames

Frame allocation involves assigning frames to pages in main memory. The challenges of frame allocation include:

- **Frame fragmentation:** When free frames are broken into small, non-contiguous blocks.
- **Frame allocation algorithms:** Algorithms such as **best fit**, **worst fit**, and **first fit** are used to assign frames to pages.

Thrashing and Virtual Memory

Thrashing occurs when a program is constantly page faulting, causing the system to spend more time retrieving pages from disk storage than executing instructions. **Virtual memory** can help to reduce thrashing by providing a larger address space, which can reduce the number of page faults.

Virtual Memory in Windows

Virtual memory in Windows involves using a combination of main memory and disk storage to provide a large address space. Windows uses a **page file** to store pages that are not in main memory.

Memory Management in Linux

Memory management in Linux involves using a combination of main memory and disk storage to provide a large address space. Linux uses a **swap space** to store pages that are not in main memory.

Memory Management in Real-Time Systems

Memory management in real-time systems involves providing predictable and fast memory access times. Real-time systems use **static memory allocation** to allocate memory at compile time, rather than at runtime.

Memory Management in Embedded Systems

Memory management in embedded systems involves providing efficient and reliable memory access. Embedded systems use **static memory allocation** and **memory-mapped I/O** to provide fast and predictable memory access times.

Case Studies of Memory Management

Case studies of memory management involve analyzing the memory management techniques used in different systems and applications. These case studies can provide insights into the trade-offs between different memory management techniques and the importance of memory management in system design.

Storage Management and File Systems

Storage Management and File Systems

Introduction to Storage Management

Storage management refers to the process of managing and optimizing the use of storage devices and media in a computer system. **Storage management** is crucial for efficient data storage and retrieval. The goals of storage management include:

- Maximizing storage capacity
- Minimizing storage costs
- Ensuring data reliability and availability
- Optimizing data access and retrieval times

File System Concept

A **file system** is a way of organizing and storing files on a storage device. It provides a hierarchical structure for storing and retrieving files. The components of a file system include:

- **Files**: collections of data stored on a device
- **Directories**: folders that contain files and other directories
- **Inodes**: data structures that contain file metadata
- **Blocks**: smallest units of storage allocation

System Calls for File Operations

System calls are APIs that allow programs to interact with the operating system. System calls for file operations include:

1. **Create**: creates a new file
2. **Delete**: deletes a file
3. **Open**: opens a file for reading or writing
4. **Close**: closes a file
5. **Read**: reads data from a file
6. **Write**: writes data to a file

Access Methods and Directory Structure

Access methods refer to the way data is accessed on a storage device. Common access methods include:

- **Sequential access**: data is accessed in a sequential manner
- **Random access**: data is accessed directly
- **Direct access**: data is accessed using a direct address The **directory structure** refers to the organization of directories and files on a storage device. Common directory structures include:

- **Hierarchical:** directories are organized in a tree-like structure
- **Flat:** all files are stored in a single directory

Disk Structure and File System Mounting

A **disk** is a storage device that consists of a series of **tracks** and **sectors**. **File system mounting** refers to the process of making a file system available for use by the operating system. The steps involved in mounting a file system include:

1. **Device identification:** identifying the storage device
2. **File system identification:** identifying the file system type
3. **Mount point creation:** creating a mount point for the file system

File Sharing and File System Implementation

File sharing refers to the ability to share files between multiple users or systems. **File system implementation** refers to the process of designing and implementing a file system. Common file system implementations include:

- **Local file systems:** files are stored on a local device
- **Network file systems:** files are stored on a remote device
- **Distributed file systems:** files are stored on multiple devices

File System Structure and Implementation

The **file system structure** refers to the organization of files and directories on a storage device. The **implementation** of a file system refers to the design and implementation of the file system. Common file system structures include:

- **Single-level directory:** all files are stored in a single directory
- **Multi-level directory:** files are stored in a hierarchical directory structure

Directory Implementation and Allocation Methods

Directory implementation refers to the design and implementation of directories on a storage device. **Allocation methods** refer to the way storage space is allocated to files and directories. Common allocation methods include:

- **Contiguous allocation:** files are stored in contiguous blocks
- **Linked allocation:** files are stored in linked blocks
- **Indexed allocation:** files are stored using an index

Free-Space Management and Efficiency

Free-space management refers to the process of managing free space on a storage device.

Efficiency refers to the optimized use of storage space. Techniques for free-space management include:

- **Bit mapping:** using a bitmap to track free space
- **Free-space lists:** maintaining a list of free space

Performance and Overview of Mass Storage Structure

Performance refers to the speed and efficiency of a storage device. The **mass storage structure** refers to the organization of storage devices in a system. Common mass storage structures include:

- **Disk arrays:** multiple disks are used to store data
- **Tape libraries:** multiple tapes are used to store data

Protection and Security in File Systems

Protection refers to the mechanisms used to protect files and directories from unauthorized access. **Security** refers to the mechanisms used to prevent unauthorized access to a system. Common protection mechanisms include:

- **Access control lists:** lists of users and their access rights
- **File permissions:** permissions assigned to files and directories

Access Control and Revocation of Access Rights

Access control refers to the mechanisms used to control access to files and directories.

Revocation of access rights refers to the process of revoking access rights to a file or directory. Techniques for access control include:

- **Discretionary access control:** access rights are assigned based on user identity
- **Mandatory access control:** access rights are assigned based on security labels

Capability-Based Systems and Language-Based Protection

Capability-based systems refer to systems that use capabilities to control access to resources. **Language-based protection** refers to the use of programming languages to enforce protection mechanisms. Common capability-based systems include:

- **Capability-based file systems:** files are protected using capabilities
- **Language-based protection mechanisms:** protection mechanisms are enforced using programming languages

File System Case Studies

Case studies refer to the analysis of real-world file systems. Common file system case studies include:

- **Unix file system:** a case study of the Unix file system
- **NTFS file system:** a case study of the NTFS file system

Future of File Systems and Storage Management

The **future of file systems** refers to the emerging trends and technologies in file systems. The **future of storage management** refers to the emerging trends and technologies in storage management. Common emerging trends include:

- **Cloud storage:** storage devices are located in the cloud
- **Object storage:** files are stored as objects
- **Distributed file systems:** files are stored on multiple devices